

SYED TAUFEEQ MUSTAFA

Software Engineer | AI, NLP, RAG, GenAI, LLM, Fine-tuning

github.com/taufeeq-mustafa | [LinkedIn](#) | +92 3352669000 | taufeeq.mustafa@gmail.com

PROFESSIONAL SUMMARY

AI/ML Engineer with hands-on experience shipping production-ready intelligent systems in **Python, FastAPI, and Django**. Skilled in **LLM integration (Anthropic, OpenAi), fine-tuning, RAG pipelines, and vector embeddings with LangChain** on **GCP Vertex AI**. Strong backend foundation on **AWS (S3, SES, RDS)** with **Redis** caching, secure API design, and **Salesforce** integration; proven delivery of multi-tenant SaaS platforms with AI-powered insights engines.

EDUCATION

SIR SYED UNIVERSITY OF ENGINEERING AND TECHNOLOGY

Karachi, Pakistan

Bachelor of Science in Software Engineering

Sep 2021 – Aug 2025

Relevant Coursework: Software Requirement Engineering, Computer Organization and Architecture, Data Structure and Algorithms, Database Systems, Operating Systems, Software Design and Architecture, Computer Communication and Networks, Software Construction and Development, Software Quality Engineering, Artificial Intelligence.

WORK EXPERIENCE

1PERCENTLABS

Karachi, Pakistan

Backend AI Engineer

Jan 2026 – Present

- Integrated Anthropic Claude (claude-opus-4-6) with web search tooling and Crunchbase API enrichment to build the PRISM AI scoring engine, evaluating services across Technical Feasibility, Risk, Budget, and Timeframe dimensions; implemented a 7-day Redis cache layer to reduce redundant LLM calls by 50% and API costs.
- Designed an invoice management system with AWS S3 presigned URL upload/download flows and a client-facing dispute workflow featuring a status lifecycle, with AWS SES email notifications to internal consultants.
- Architected a Salesforce-triggered account merge system with a dual-path engine, integrating Apex triggers with Django REST endpoints via API-key authentication and enforcing validation guards for countersigned applications, in-flight Stripe payments, and idempotency checks.

WENAWA

Karachi, Pakistan

Artificial Intelligence Trainee

Oct 2025 – Dec 2025

- Packaged a **LangChain-based RAG pipeline** (Llama-3 + FAISS vector store) into a **FastAPI + Docker** microservice handling **~500 daily queries** with avg response time under **1.5s**.
- Built a Python pipeline (**pandas, spaCy, NLTK**) that ingested **~50K** raw customer reviews, applied tokenization, stop-word removal, and lemmatization to produce a labeled dataset for sentiment-model training.
- Performed exploratory data analysis on **100K+ records** in **pandas/matplotlib/seaborn**, surfaced class imbalances and data-quality issues, and delivered Git-tracked reproducible notebooks to the senior ML team.

UNIVERSITY PROJECTS

NLP-Based Talking Therapy Platform (FYP)

Aug 2025

- Fine-tuned facebook/blenderbot-400M-distill on a synthetic dataset for domain-specific responses and local hosting.
- Implemented FastAPI for efficient real-time API communication **and integrated LangChain vector store embeddings** for context-aware multi-session continuity.

HunarShehar (Google APAC Hackathon)

May 2025

- Deployed AI agentic model on Google Vertex AI and integrated Google Agent Development Kit for scalable, personalized onboarding with adaptive assessments, enabling data-driven user profiling.
- Built AI agents using Gemini for bilingual microlessons and role-play simulations, enhancing engagement by 30%.

SKILLS AND CERTIFICATIONS

Technical Skills: Python, Django, SQL, FastAPI, Node.js, Docker, GCP, LangChain, RAG, AI Agents, AWS, Jupyter, Git.

Certifications: AI for Devs Certification from DataCamp, Object Oriented Programming (Coursera - UC San Diego).